

# Analyzing Simulation Output: Batch Means Intervals

DCS 307: Simulation  
Winter 2023

29 March 2023

# Great Ideas in Simulation

- ① Lehmer RNGs (linear congruential generators)

$$x_{i+1} = ax_i \bmod m$$

- ② Probability integral transformation:

$$F(X) \sim U(0, 1)$$

- ③ Algorithms for variate generation: inversion, acceptance/rejection, composition, special properties
- ④ Handling autocorrelation in constructing confidence intervals
- ⑤ Next-event simulation: e.g., use arrival to create next arrival
- ⑥ Agent-based simulation

# Great Ideas in Simulation

- ① Lehmer RNGs (linear congruential generators)

$$x_{i+1} = ax_i \bmod m$$

- ② Probability integral transformation:


$$F(X) \sim U(0, 1)$$

- ③ Algorithms for variate generation: inversion, acceptance/rejection, composition, special properties
- ④ Handling autocorrelation in constructing confidence intervals
- ⑤ Next-event simulation: e.g., use arrival to create next arrival
- ⑥ Agent-based simulation

# Recap: Central Limit Theorem

- Choose *any* of the distributions encountered so far
- Generate a sample of  $n > 1$  variates (mean  $\mu$ , sd  $\sigma$ )
- Generate  $j > 1$   $n$ -point samples, compute  $\bar{x}$  for each

$$\underbrace{x_1, x_2, \dots, x_n}_{\bar{x}_1}, \underbrace{x_{n+1}, x_{n+2}, \dots, x_{2n}}_{\bar{x}_2}, \underbrace{x_{2n+1}, x_{2n+2}, \dots, x_{3n}}_{\bar{x}_3}, x_{3n+1}, \dots$$



- For large  $n$ , histogram density approximates *normal* $(\mu, \sigma/\sqrt{n})$

## Recap (Cont.)

- CLT: mean of *iid* sequence of RVs tends to *normal*

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i \rightarrow \text{normal}(\mu, \sigma/\sqrt{n})$$

- The  $x_i$  correspond to simulation estimates, i.e., samples
  - e.g., average wait, average queue length, etc.

- We want the (unknown)  $\mu$  from *normal*( $\mu, \sigma/\sqrt{n}$ )

- $\bar{x}$  is an estimate of  $\mu$

- Standardized sample means tend to *normal*(0,1)
- Using t-statistic, approximates *Student*( $n-1$ )

$$t = \frac{\bar{x}_j - \mu}{s/\sqrt{n-1}}$$

## Recap (Cont.)

- CLT: mean of *iid* sequence of RVs tends to *normal*

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i \rightarrow \text{normal}(\mu, \sigma/\sqrt{n})$$

- The  $x_i$  correspond to simulation estimates, i.e., samples
  - e.g., average wait, average queue length, etc.
- We want the (unknown)  $\mu$  from *normal*( $\mu, \sigma/\sqrt{n}$ )
  - $\bar{x}$  is an estimate of  $\mu$
- Standardized sample means tend to *normal*(0,1)
- Using t-statistic, approximates *Student*( $n-1$ )

$$t = \frac{\bar{x}_j - \mu}{s/\sqrt{n-1}}$$

## Recap (Cont.)

- CLT: mean of *iid* sequence of RVs tends to *normal*

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i \rightarrow \text{normal}(\mu, \sigma/\sqrt{n})$$

- The  $x_i$  correspond to simulation estimates, i.e., samples
  - e.g., average wait, average queue length, etc.
- We want the (unknown)  $\mu$  from *normal*( $\mu, \sigma/\sqrt{n}$ )
  - $\bar{x}$  is an estimate of  $\mu$
- Standardized sample means tend to *normal*(0,1)
- Using t-statistic, approximates *Student*( $n-1$ )

$$t = \frac{\bar{x}_j - \mu}{s/\sqrt{n-1}}$$

# Interval Estimation Algorithm

## Algorithm

To calculate an *interval estimate* for the unknown mean  $\mu$  of the *population* from which a random sample  $x_1, x_2, x_3, \dots, x_n$  was drawn:

- Pick a level of confidence  $1 - \alpha$ 
  - typically  $\alpha = 0.05$  (95% confidence)
- Calculate the sample mean  $\bar{x}$  and standard deviation  $s$
- Calculate the critical value  $t^* = \text{qt}(1 - \alpha/2, n - 1)$
- Calculate the interval endpoints


$$\bar{x} \pm \frac{t^* s}{\sqrt{n - 1}}$$

If  $n$  is sufficiently large, then you are  $(1 - \alpha) \times 100\%$  confident that the mean  $\mu$  lies within the interval. The midpoint of the interval is  $\bar{x}$ .



# Example

- 100 samples of size  $n = 9$  drawn from  $normal(6, 3)$  population
- For each sample, construct an exact 95% confidence interval
- 94 intervals cover  $\mu = 6$
- 3 missed high (red), 3 missed low (red)



# Producing Confidence Intervals in R: Discrete-Event

- Recall that `ssq()` service times are *exponential* ( $\lambda = 10/9$ )
- Use `t.test()` approach to produce 95% confidence intervals

$$\bar{x} \pm t_{\alpha/2, n-1} \frac{s}{\sqrt{n-1}}$$


---

```
set.seed(12345)
seeds <- sample(10^6:(10^7 - 1), size = 100, replace = FALSE)

ci_lo <- numeric(100) # vector for 100 CI low endpoints
ci_hi <- numeric(100) # vector for 100 CI high endpoints

numJobs <- 30
for (i in 1:100) {
  output <- ssq(numJobs, seeds[i], showOutput = FALSE,
                saveServiceTimes = TRUE)
  ci <- t.test(output$serviceTimes, conf.level = 0.95)
  ci_lo[i] <- ci$conf.int[1]
  ci_hi[i] <- ci$conf.int[2]
}
```

# Confidence Intervals for Discrete-Event Simulation



- 92 intervals cover  $\mu = 0.9$
- 8 intervals missed low (red)
- CIs are approximate when population distribution not *normal*

# Confidence Intervals for Discrete-Event Simulation


- The distribution of wait times for ssq is unknown
  - Try `t.test()` approach to produce 95% confidence intervals
- 

```
ci_lo <- rep(NA, 100) # vector for 100 CI low endpoints
ci_hi <- rep(NA, 100) # vector for 100 CI high endpoints

numJobs <- 30
warmup <- 1000
for (i in 1:100) {
  output <- ssq(numJobs + warmup, seeds[i], showOutput = FALSE,
               saveWaitTimes = TRUE)
  waits <- output$waitTimes[-(1:warmup)]

  ci <- t.test(waits, conf.level = 0.95)
  ci_lo[i] <- ci$conf.int[1]
  ci_hi[i] <- ci$conf.int[2]
}
```


# Confidence Intervals for Discrete-Event Simulation



- For an  $M/M/1$  queue with arrival rate  $\lambda$  and service rate  $\nu$ , the theoretical steady-state wait is  $1/(\nu - \lambda) - 1/\nu = 8.1$
- Despite claim of 95% confidence, only 10 of 100 CIs bracket
- 28 miss high, 62 miss low

# Effect of Initial Conditions: “Warm Up”

- ssq: accumulated  $\bar{w}$  printed every 100 jobs




- Again, the theoretical steady-state wait is 8.1

# Queueing Wait Times



- How does a long wait for the first person in line affect...
  - second in line? third in line? ... far back in line?

# Example: Effect of Autocorrelation




(Note: different parameter values here than in other examples)



# Autocorrelation of ssq Wait Times

- In ssq, strong positive autocorrelation b/w wait times




---

```
output <- ssq(10000, 1234567, showOutput = F, saveWaitTimes = T)
waits <- output$waitTimes[-(1:1000)] # remove "warm-up"
acf(waits)
```

# Autocorrelation of ssq Wait Times

- Strong positive autocorrelation even for large lags



---

```
output <- ssq(10000, 1234567, showOutput = F, saveWaitTimes = T)
waits <- output$waitTimes[-(1:1000)] # remove "warm-up"
acf(waits, lag.max = 300)
```


# Confidence Intervals in Presence of Autocorrelation

- Method of Batch Means:
  - (a) Make *one* long simulation run, deleting warmup
  - (b) Partition into  $n$  batches each of batch size  $b$
  - (c) Compute a sample mean for each batch
  - (d) Construct an interval estimate using means of batches

$$\underbrace{x_1, x_2, \dots, x_b}_{\bar{x}_1}, \underbrace{x_{b+1}, x_{b+2}, \dots, x_{2b}}_{\bar{x}_2}, \dots, \underbrace{x_{(n-1)b+1}, x_{(n-1)b+2}, \dots, x_{nb}}_{\bar{x}_n}$$


- Each batch: like a “transient simulation” w/o initial state bias
- State @ beginning of each batch is state @ end of previous

# Batch Means for ssq Wait Times



- 100 different 95% confidence intervals using 100 diff. seeds
- Number of batches  $n = 32$ , batch size  $b = 2048$
- First 1024 jobs discarded to mitigate start-up bias
- Batch means reduces the impact of autocorrelated data

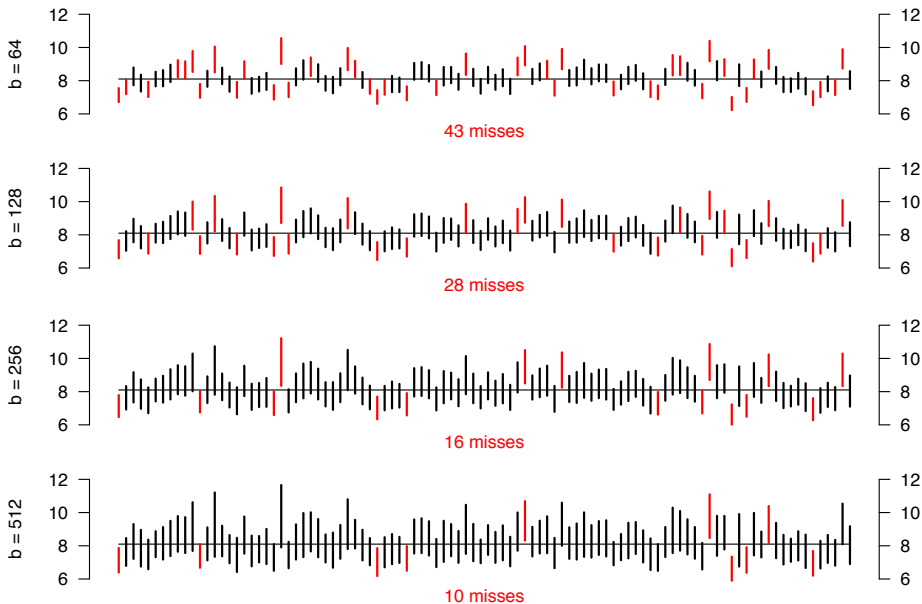
# Batch Means for ssq Wait Times



```
numJobs <- 65536; warmup <- 1024
b <- 2048; n <- numJobs / b

output <- ssq(numJobs + warmup, 1234567, showOutput = FALSE, saveWaitTimes = TRUE)
waits <- output$waitTimes[-(1:warmup)] # remove warmup
batchMeans <- numeric(n)
for (j in 1:n) {
  batch <- waits[(1:b) + (b * (j - 1))]
  batchMeans[j] <- mean(batch)
}
ci <- t.test(batchMeans, conf.level = 0.95)
```

# Batch Size & Confidence: ssq with $b = 64, 128, 256, 512$



# Effect of Batch Parameters

- Consider any one confidence interval from previous figures
- Choice of  $(b, n)$  has *no* impact on the *point* estimate
- Only the width of the 95% confidence interval is affected
- Use `ssq`<sup>1</sup> to generate 65536 waits with  $(\lambda, \nu) = (1, 10/9)$
- Batch means with different  $(b, n)$ :

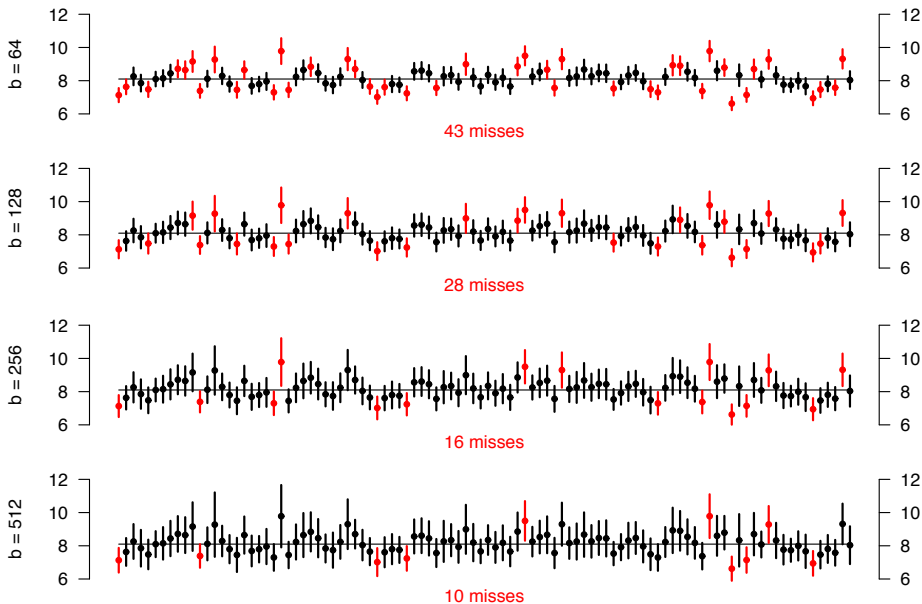
$(b, n)$	(32, 2048)	(128, 512)	(512, 128)	(2048, 32)
$\bar{w}$	$8.68 \pm 0.40$	$8.68 \pm 0.73$	$8.68 \pm 1.14$	$8.68 \pm 1.38$

- Note that  $\bar{w} = 8.68$  is independent of  $(b, n)$
- Width of the interval estimate is not

---

<sup>1</sup>Use `seed = 1234567`, `maxArrivals = 65536 + 1024`; discard first 1024 waits.

# Batch Size & Confidence: ssq with $b = 64, 128, 256, 512$






# Choosing Batch Size

- Schmeiser (1982) suggests number of batches  $k$  between 10 and 30
- Pegden et al. (1995) recommend a batch size  $b$  that is at least  $10\times$  as large as the largest lag for which autocorrelation remains significant
- Banks et al. (2005) recommend  $b$  be increased until the lag-one autocorrelation between batch means is less than 0.2

# In-class Work: Meaning of Confidence



- Create figures like above, but using craps instead of *normal*:
  - Fig 1: 100 CIs with each CI via 9 samples from craps(10)
  - Fig 2: 100 CIs with each CI via 9 samples from craps(100)
- For each sample, construct a 95% CI using t.test
- Color intervals that miss using red (true prob: 244/495)

# In-class Work: Getting You Started

- Create an empty plot with true theoretical value:

```
plot(NA, NA, xlim = c(1,100), ylim = c(ylo,yhi),  
     bty = "n", xaxt = "n",  
     xlab = "", ylab = "", las = 1)  
exact <- 244/495  
abline(h = exact)
```

- Use segments to draw each CI
- To create a 9-point vector of craps estimates:

```
values <- sapply(1:9,  
  function(i) { craps(nreps, showProgress=FALSE) } )
```

- ① Repeat the previous batch-means experiments using sojourn times
- ② For each of  $b = 64, 128, 256, 512, 1024, 2048$ :
  - Produce a figure of 100 different 95% confidence intervals
  - Use a different initial seed for producing each CI
  - Use red for intervals that miss, black for intervals that do not
  - In the figure, indicate the batch size and number of misses
  - Theoretical  $\bar{o}$  for default ssq is  $1/(\nu - \lambda) = 9.0$

(Note: each of these figures will require non-trivial execution time)